
Object Recognition : An Exploration

Anubhab Majumdar
amajumd@ncsu.edu

Toshal Phene
tphene@ncsu.edu

Shubham Munot
samunot@ncsu.edu

1 Introduction

Object recognition is a quintessential machine learning problem of modern times. It is difficult for machines to recognize a large variety of objects in different conditions of lighting, occlusion and skew. In recent years, this field has seen huge leaps with machines becoming adept at object recognition. In this project, we explore different approaches of tackling this problem

We are using the CALTECH 256 dataset ([link](#)) for object recognition [1]. The dataset has 256 unique categories of images. The number of samples in each category varies from 80 to over 200. This dataset is quite a standard dataset in the field of object recognition and has been used in important research papers like [2]. We experimented on a small subset (4-5 categories) of the dataset.

The experiments have been implemented using student version of MATLAB and python/tensorflow. Some of the less computationally intensive models are trained on personal computers, while the deep networks are trained on VCL servers.

2 Methodology

A typical classification system is shown in Fig.1. The paper focuses on the machine learning block shown in Fig.1. Different learning algorithms ranging from simple and lazy methods like KNN to deep convolution neural network is applied and results are compared for a comprehensive understanding of which method works better. The paper also tries to draw conclusion as to why certain methods work better than others.




Figure 1: A typical prediction system

Before jumping into the algorithms applied, let's talk about the data itself. We have chosen the following 4 categories for our experimentation:



Figure 2: Few samples from our dataset

- Backpack (Sample size - 151)
- Binoculars (Sample size - 216)
- Eiffel Tower (Sample size - 83)
- Fried egg (Sample size - 90)

Sample of some images from the dataset are shown in Fig.2.

2.1 Baseline

For determining baseline, we have used the simple technique of majority class. As seen from above, among the 4 different classes we are experimenting with, binocular has the largest number of sample images. Therefore, given any image from the dataset, the probability of that image having a binocular is highest among the 4 classes. And this is what we use as baseline method.

2.2 KNN

Gabor filters are good models of how humans distinguish texture, and are therefore a useful model to use when designing algorithms to recognize texture. This example uses the basic approach described in (A. K. Jain and F. Farrokhnia, "Unsupervised Texture Segmentation Using Gabor Filters",1991) to perform texture segmentation.[4]

Design an array of Gabor Filters which are tuned to different frequencies and orientations. The set of frequencies and orientations is designed to localize different, roughly orthogonal, subsets of frequency and orientation information in the input image. Regularly sample orientations between [0,150] degrees in steps of 30 degrees. Sample wavelength in increasing powers of two starting from $4/\sqrt{2}$ up to the hypotenuse length of the input image. These combinations of frequency and orientation are taken from [Jain,1991] cited in the introduction.[4]

Extract gabor magnitude features from source image. When working with Gabor filters, it is common to work with the magnitude response of each filter. Gabor magnitude response is also sometimes referred to as "Gabor Energy". Each MxN Gabor magnitude output image in `gabormag(:, :, ind)` is the output of the corresponding gabor filter `g(ind)`. [4]

A nearest-neighbor classification object, where both distance metric ("nearest") and number of neighbors can be altered. The object classifies new observations using the predict method. The object contains the data used for training, so can compute resubstitution predictions.

2.3 Convolutional Neural Network

Fully connected neural networks (NN) are not the most suitable architecture to classify images because such a network architecture does not take into account the spatial structure of the images.

The spatial location information of individual pixels are lost as soon as we transform it's 2D structure into one dimension to input into the neural network architecture. To solve this problem, LeCun et. al. came up with convolutional neural network (CNN) in their pioneering work LeNet-5 [7].

CNN is a type of feed forward artificial neural network that retains the location information of training data and is well adapted to classify images. The connectivity pattern between it's neurons is inspired by the organization of the animal visual cortex. A typical CNN architecture is shown in Fig.3.




Figure 3: A typical CNN architecture

CNN has 3 components:

- Convolutional layer - These layers use local receptive field and a set of filters to convolve across the width and height of the input volume, computing the dot product between the entries of the filter and the input and producing a 2-dimensional activation map of that filter. As a result, the network learns filters that activate when they see some specific type of feature at some spatial position in the input.
- Pooling layer - These layers immediately follow convolutional layers. What the pooling layers do is simplify the information in the output from the convolutional layer. In our experiments we have used max-pooling.
- Fully connected layer - These layers are identical to hidden layers of an ANN. They connect the last pooling layer with the output layer. The last layer produces the classification results.

However, building and training a deep CNN architecture (ConvNet) is not easy because:

- There is no exact set of rules which determine how deep or how wide the network should be to produce good accuracy
- Small number of data samples per class in Caltech-256 dataset
- Huge computational cost of training a ConvNet

There are 4 main aspects of ConvNet that we need to determine:

- Training data
 - How large the training set should be? Should the training set size be increased by augmenting the available images?
 - What should be the resolution of each input image?
 - What should be the ideal batch size for training the network?
- Convolutional layers
 - What should be the depth of the network, i.e., how many convolutional layers should be present?
 - What should be the width of the network, i.e., how many features should be present per convolutional layer?
- Fully connected layer
 - What should be the depth of the fully connected layer?
 - What should be the size of each fully connected layer, i.e. how many neurons should be present in each fully connected layer?
- Prevent overfitting - CNN, like artificial NN can easily overfit. How to prevent that?

Our focus in this paper is more on trying to understand how tweaking each of the above mentioned aspect affects the performance of ConvNet, rather than to try everything to somehow achieve good accuracy.

Since the number of samples per class is quite less, we have increased our training set using image augmentation. We kept 80 images aside for testing. The training set contains 3680 images - 460 are unique images and rest are augmented using rotation of image and adding Gaussian blur.

We have used dropout, a regularization technique, to prevent overfitting. It is proposed by Hinton et. al. in their seminal paper [8]. The term "dropout" refers to dropping out units (both hidden and visible) in a neural network to prevent overfitting.

The networks are created with python and tensorflow. The codebase is available in GitHub [9].

3 Experimental Results

3.1 Baseline

	Backpack	Binocular	Eiffel Tower	Fried Egg
Backpack	0	20	0	0
Binocular	0	20	0	0
Eiffel Tower	0	20	0	0
Fried Egg	0	20	0	0

3.2 KNN

	Backpack	Binocular	Eiffel Tower	Fried Egg
Backpack	12	5	0	3
Binocular	3	17	0	0
Eiffel Tower	3	5	9	3
Fried Egg	3	8	1	8

3.3 Convolutional Neural Network

Index	Input image size	Conv layer 1	Conv layer 2	Conv layer 3	Fully connected layer 1	Fully connected layer 2	Train accuracy	Test accuracy
CNN 1	28x28x1	14x14x32	7x7x64	NA	64x1	NA	45.43%	26.25%
CNN 2	28x28x1	14x14x32	7x7x64	NA	1024x1	NA	79.78%	27.50%
CNN 3	64x64x1	32x32x32	16x16x64	8x8x128	1024x1	NA	91.00%	42.50%
CNN 4	28x28x1	14x14x64	7x7x128	NA	1024x1	1024x1	96.82%	69.69%
CNN 5	28x28x1	14x14x128	7x7x256	NA	1024x1	1024x1	96.54%	67.81%

Table 1: Results of 5 different CNN architectures

Some interesting observations from the result above are:

- A very shallow network cannot capture enough information about the different categories of images to provide good testing accuracy (CNN 1 and CNN 2)
- Increasing the depth of convolution layer results better testing accuracy (CNN 3). This can be attributed to the fact that more complicated features are captured by the convolutional layers, which describes the objects more accurately
- Increasing depth of fully connected layers, even with less convolutional layers, boosts accuracy by a large percentage (CNN 4). We believe this is because more hidden layers learn hierarchical representations of the input data and can better generalize to unseen data as well

- We do not achieve better accuracy by just making the network wide and not deep (CNN 5). The reason, we believe, is even though there are more features, they are not much different from the existing ones to help distinguish the categories any better than CNN 4

4 Conclusion

The project was an exploration across various approaches to object recognition across time. We started with simpler model like KNN and moved to extremely complex models like CNN and everything in between. Fig.4 shows the comparison of the various experimental results we captured in this project.

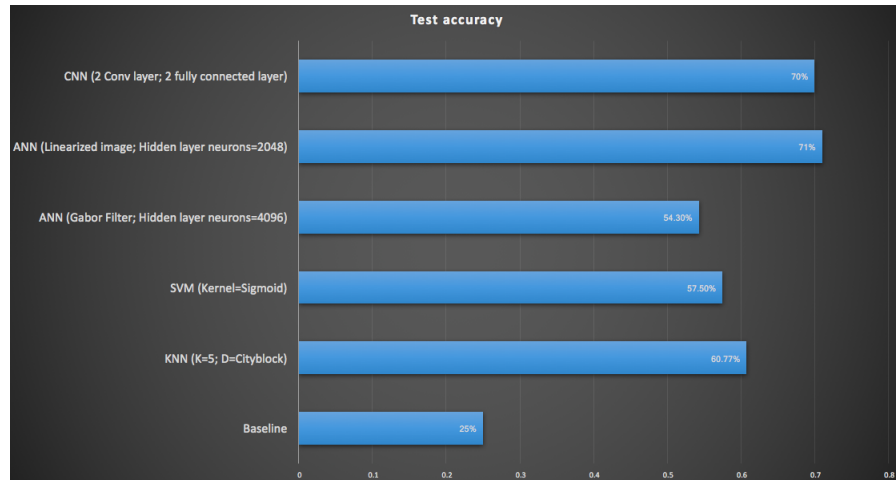


Figure 4: A comparison of experimental results

The comparison of the results yield some very interesting points:

- Neural networks trained with the raw image itself (i.e. without any feature engineering) produces much better results in terms of accuracy than algorithms with feature engineering.
- Feature engineering for large categories of object is fragile. There are lot of parameters that need to be decided to get good accuracy and there is no guarantee that introduction of a new category won't degrade the performance.
- One of the best advantage of deep convolutional neural network is the independence from feature engineering. This makes the approach more robust.

The take home point, for us, from doing this project are:

- Handcrafted features for large variety of images is unreliable and fragile
- Featureless approach to computer vision problem with deep neural network yields better accuracy
- Modern computer vision research is moving towards featureless approach with deep ConvNet. Ex: ImageNet and GoogLeNet

References

- [1] Greg Griffin, Alex Holub and Pietro Perona.
Caltech-256 Object Category Dataset.
- [2] Hao Zhang, Alexander C. Berg, Michael Maire and Jitendra Malik
SVM-KNN: Discriminative Nearest Neighbor Classification for Visual Category Recognition.
- [3] Knuth: Computers and Typesetting,
<http://www-cs-faculty.stanford.edu/~uno/abcde.html>

- [4] Texture Segmentation using Gabor Filters
<https://www.mathworks.com/help/images/texture-segmentation-using-gabor-filters.htm>
- [5] Simona E.Grigorescu, Nicolai Petkov, and Peter Kruizinga
Comparison of Texture Features based on Gabor Filters.
- [6] Anil K. Jain,Farshid Farrokhnia
Unsupervised Texture Segmentation Using Gabor Filters.
- [7] Yann Lecun, Leon Bottou, Yoshua Bengio and Patrick Haffner
Gradient-Based learning applied to document recognition
- [8] Nitish Srivastava, Geoffrey Hinton, Alex Krizhevsky, Ilya Sutskever and Ruslan Salakhutdinov
Dropout: A Simple Way to Prevent Neural Networks from Overfitting
- [9] <https://github.com/anubhabMajumdar/Object-Recognition-using-ConvNet>